

Current Recursion and Tensor-Network Execution in PYAMPLICOL

Purpose. This note summarizes the matrix-element generation strategy used by AMPLICOL and the corresponding PYAMPLICOL implementation. The common idea is to avoid explicit Feynman-diagram enumeration. Instead, matrix elements are assembled from a directed acyclic graph (DAG) of off-shell currents. AMPLICOL emits Fortran code that numerically sweeps this current table. PYAMPLICOL builds the same logical DAG in Python, lowers bounded pieces to tensor-network expressions handled by Spenso, compiles the resulting parametric scalar/vector expressions with Symbolica evaluators, and can now execute the staged sweep from Rust through the PyO3 runtime `rusticol`.

1 The AmpliCol Current Recursion

For a fixed leading-color ordering, a current is labelled by the external subset it contains, the particle type carried by the off-shell line, and any spin/chirality information. Schematically, if I is a set or ordered interval of external colored legs, a current is computed as

$$\mathcal{J}_{t,\chi}^\alpha(I) = P^\alpha_\beta(t, p_I) \sum_{\substack{I=A \cup B \\ A \cap B = \emptyset}} \mathcal{V}^\beta_{\gamma\delta}(t_A, t_B \rightarrow t) \mathcal{J}_{t_A, \chi_A}^\gamma(A) \mathcal{J}_{t_B, \chi_B}^\delta(B). \quad (1)$$

Here t is the current type (gluon, quark, auxiliary tensor, etc.), χ is the Weyl chirality where relevant, $p_I = \sum_{i \in I} p_i$, P is the propagator, and $\mathcal{V}$ is the appropriate Feynman-rule tensor. The four-gluon vertex is treated through the auxiliary antisymmetric tensor route, so it is still represented by cubic current combinations.

AMPLICOL's optimized mode does not loop over full helicity configurations as independent computations. It weaves helicity information into the current table itself. Each external helicity state is a source current with a unique bit in an `ext_cur`-style bitset. An internal current is identified by

$$K(\mathcal{J}) = (t, \chi, \text{bin}(I), \text{ext_cur}), \quad (2)$$

where `bin(I)` records the external subset and `ext_cur` is the union of the source-current bits feeding the current. When a new vertex contribution produces the same key as an existing current, AMPLICOL appends that vertex to the existing current rather than creating a duplicate node. After all currents are constructed, `curr2amp` records the two endpoint currents whose contraction gives each retained helicity amplitude.

The generated Fortran library has the same layered structure:

```
subroutine evaluate_amp(p, amps)
  complex(kind=8) :: val_c(1:6, n_cur)
  complex(kind=8) :: int_c(1:6, n_vert)

  call compute_external_currents(pp, val_c)
  do size = 2, n_external_minus_one
    call compute_vertices_size(pp, val_c, int_c)
    call compute_currents_size(pp, val_c, int_c)
  end do
  call compute_amps(amps, val_c)
end subroutine
```

The arrays `val_c` and `int_c` are the numerical current and interaction tables. This is the performance-critical idea that PYAMPLICOL keeps: one topological sweep produces all retained amplitudes, and shared parent currents are evaluated once.

2 The pyAmpliCol DAG

PYAMPLICOL first builds an explicit Python representation of the same current table. In the generic compiler, a current is identified by a model-driven physics index rather than by a process-family name. The identity used for deduplication is

$$K_{\text{PYAMPLICOL}} = (\text{particle_id}, L, H, \chi, S, F, Q, C, P, A), \quad (3)$$

where L is the sorted tuple of external labels, H is the AMPLICOL-style helicity-source ancestry bitset, χ is chirality, S is any spin-state tag needed by the source or propagator, F is flavour flow, Q is charge flow, C is the leading-color

state, P is the momentum mask, and A is an optional auxiliary kind such as the antisymmetric tensor used for the four-gluon vertex. The implementation also carries `ordered_external_labels`. That field is leading-color ordering metadata used by the color engine and amplitude grouping, but it is deliberately not part of equality or hashing. Deduplication is therefore exactly equality of Eq. (3); no rule is allowed to ask whether the process is “ Z plus jets” or any other family.

field	role in the current identity
<code>particle_id</code>	Model/UFO particle id, e.g. 21, 23, 1, -1 , or an auxiliary id.
<code>external_labels</code> and <code>external_mask</code>	Which external legs are contained in the current. The mask is the bit representation of the labels.
<code>helicity_ancestry</code>	Bitwise union of the source-current helicity bits feeding this current. This is how helicity configurations are woven into one shared DAG.
<code>chirality</code> and <code>spin_state</code>	Local spin information, especially for Weyl fermions and massive-vector sources.
<code>flavour_flow</code> and <code>charge_flow</code>	Quantum-number flow checked locally by model vertices, e.g. $u \rightarrow dW^+$ or $e^+e^- \rightarrow \gamma/Z$.
<code>color_state</code>	Leading-color open-line or trace state now; later the hook for Idenso-backed NLC/full-color basis data.
<code>momentum_mask</code>	Which all-outgoing momenta are summed for propagators and momentum-dependent vertices.
<code>auxiliary_kind</code>	Optional tag distinguishing auxiliary tensor currents from ordinary particles.

Source currents are built from external leg metadata. Internal currents are then generated by increasing current size, with all interactions that feed the same current index attached to one node.

```

for external helicity source s:
    add_current(key=(pdg_s, labels_s, chirality_s, bit_s), is_source=True)

for current_size in increasing_sizes:
    for compatible split A | B:
        result_key = combine_key(key[A], key[B], vertex_kind)
        result = add_or_reuse_current(result_key)
        add_interaction(left=A, right=B, result=result, vertex=vertex_kind)

for retained helicity ancestry:
    add_amplitude(left=current_with_Z_and_gluons, right=anti_current)

prune_dead_currents_after_helicity_filter()

```

At runtime, PYAMPLICOL mirrors the Fortran table layout with contiguous arrays:

$$C_{s,i,c} \equiv \text{current_rows}[s, i, c], \quad T_{s,v,c} \equiv \text{interaction_rows}[s, v, c],$$

where s is the phase-space sample in the current batch, i is a current id, v is an interaction id, and c is a component index. External wavefunctions fill source-current rows once per sample. Bounded evaluator blocks then fill interactions and derived currents stage by stage:

```

fill_source_currents(current_rows, momenta, helicity_sources)

for stage in current_size_layers:
    state = ParamBuilder.pack(current_rows, interaction_rows, momenta)
    interaction_rows[:, stage.vertices] = vertex_eval[stage](state)
    current_rows[:, stage.currents] = current_eval[stage](state)

amps = amplitude_eval(ParamBuilder.pack(current_rows, momenta))
me2 = normalization * sum(weight[h] * abs(amps[h])**2 for h in retained_h)

```

This is why the evaluator-only benchmark is meaningful: the expensive recursion is inside compiled evaluator calls, while Python only performs source wavefunction construction, parameter packing, and table transfers.

3 Spenso Tensor Networks and hep_lib

Spenso is used to express Lorentz, spinor, auxiliary-tensor, and eventually color contractions in a representation-aware way. A tensor-network expression contains named tensors with typed slots:

$$\mathcal{V}_{\mu\nu}{}^\rho(p_A, p_B) A^\mu B^\nu, \quad \text{or in code: } \text{V}(\text{m}(\text{"a"}), \text{m}(\text{"b"}), \text{m}(\text{"out"})) * \text{A}(\text{m}(\text{"a"})) * \text{B}(\text{m}(\text{"b"})).$$

The expression is not expanded. It is a compact product of tensor factors with shared representation labels. The library `hep_lib` is the dictionary that tells Spenso what each tensor name means: dense model tensors, metric tensors, momentum-dependent propagators, and parametric rank-one tensors representing runtime inputs.

```
library = model.build_tensor_library()      # returns TensorLibrary.hep_lib_atom()
builder = ParamBuilder()
mink = Representation.mink(4)

builder.register_rank1_tensor(
    library,
    tensor_name="current_A",
    representation=mink,
    head=("stage", "A"),
    length=4,
    role="block_current",
)
builder.register_rank1_tensor(
    library,
    tensor_name="p_A",
    representation=mink,
    head=("stage", "pA"),
    length=4,
    role="block_momentum",
    real_valued=True,
)

raw = V3g(mink("a"), mink("b"), mink("out"), p_A="p_A", p_B="p_B") \
    * TensorName("current_A")(mink("a")).to_expression() \
    * TensorName("current_B")(mink("b")).to_expression()

network = TensorNetwork(raw, library)
network.execute(library=library)
out_tensor = network.result_tensor(library)
evaluator = out_tensor.evaluator(params=builder.parameter_symbols())
```

The important point is that `hep_lib` stays fixed during the execution. New graph currents are not registered as permanent physics tensors. Instead, the current values are runtime parameters for a bounded block. Spenso executes the factorized tensor network, contracts the representation indices, and returns a parametric output tensor. Symbolica then compiles that parametric tensor to an evaluator. For high multiplicity, PYAMPLICOL interleaves this process: after a current or current-size layer is contracted, its numerical output is written back into the current table and becomes an input parameter for later blocks. This avoids one giant scalar expression and keeps the DAG close to AMPLICOL's execution model.

If explicit color tensors are introduced in future full-color modes, color simplification belongs before Spenso execution:

```
raw ← idenso.simplify_color(raw),      then build TensorNetwork(raw, hep_lib).
```

For the current leading-color milestone, color is already reduced to scalar factors associated with the ordered amplitude.

4 Example: $d\bar{d} \rightarrow Zgg$

This section spells out one concrete DAG so that the split between compiled current contractions and Rust-interpreted scheduling is explicit. The user-level process is

$$d(p_1) \bar{d}(p_2) \rightarrow Z(p_3) g(p_4) g(p_5).$$

PYAMPLICOL stores the process in the all-outgoing convention. Thus labels $1, \dots, 5$ carry the fields

$$\bar{d}_1, \quad d_2, \quad Z_3, \quad g_4, \quad g_5,$$

with momenta $-p_1, -p_2, p_3, p_4$, and p_5 , respectively. A label set such as $\{2, 4, 5\}$ is represented by $L = (2, 4, 5)$ and by the mask

$$m(2, 4, 5) = 2^{2-1} + 2^{4-1} + 2^{5-1} = 26.$$

The momentum mask is the same bit pattern in this example, and it tells Rusticol which stored momentum sum must be passed to the next compiled evaluator.

For one leading-color sector, the colored ordered labels are

$$O = (2, 4, 5, 1),$$

with the color-singlet Z_3 allowed to attach anywhere on the quark line. The color engine may also provide the sector $O = (2, 5, 4, 1)$. Nothing in the DAG builder special-cases this as a “ $Z + 2g$ ” process: source currents are ordinary external-leg currents and all later currents are discovered from model vertices.

Source current indices

For a fixed retained source helicity choice, the source entries look as follows. The bit b_i denotes the unique helicity-source bit assigned to that external state.

source	current	mask	index payload
S_1	$\mathcal{J}_{\bar{d}}(\{1\})$	1	(pid = -1, $L = (1)$, $H = b_1, \chi, S, F = (-1), Q, C_{\bar{q}}, P = 1$)
S_2	$\mathcal{J}_d(\{2\})$	2	(pid = 1, $L = (2)$, $H = b_2, \chi, S, F = (1), Q, C_q, P = 2$)
S_3	$\epsilon_Z(\{3\})$	4	(pid = 23, $L = (3)$, $H = b_3, S = h_Z, C_{\text{singlet}}, P = 4$)
S_4	$\epsilon_g(\{4\})$	8	(pid = 21, $L = (4)$, $H = b_4, S = h_4, C_g, P = 8$)
S_5	$\epsilon_g(\{5\})$	16	(pid = 21, $L = (5)$, $H = b_5, S = h_5, C_g, P = 16$)

The actual stored index also contains the exact leading-color sector id, line-group metadata, and the current dimension. For this example the quark and anti-quark source currents have Weyl-spinor components, gluons and the Z have Lorentz-vector components, and auxiliary tensor currents have the antisymmetric tensor components required by the cubicized four-gluon rule.

What each compiled stage computes

The generic stage compiler groups interactions by the size of the result current. For each group it emits a multi-output Symbolica evaluator. That evaluator is compiled once and receives only parent-current components, momentum sums, and real couplings as parameters. It returns the components of all result currents in the stage.

For the sector $O = (2, 4, 5, 1)$, the schematic stage contents are:

$$E_2 : \quad \mathcal{J}_g(4, 5) = P_g(p_{45}) \mathcal{V}_{3g}(S_4, S_5; p_4, p_5), \quad (4)$$

$$\mathcal{J}_T(4, 5) = \mathcal{V}_{ggT}(S_4, S_5), \quad (5)$$

$$\mathcal{J}_d(2, 4) = P_q(p_{24}) \mathcal{V}_{ggq}(S_2, S_4; p_2, p_4), \quad (6)$$

$$\mathcal{J}_d(2, 3) = P_q(p_{23}) \mathcal{V}_{qZq}(S_2, S_3; p_2, p_3), \quad (7)$$

plus the analogous allowed currents for the other color sector and for the anti-quark side. The notation E_2 means a compiled evaluator for all two-label result currents, not a Python function that loops over these formulas component by component.

The next stage consumes the rows written by E_2 :

$$E_3 : \quad \mathcal{J}_d(2, 4, 5) = P_q(p_{245}) [\mathcal{V}_{ggq}(\mathcal{J}_d(2, 4), S_5) + \mathcal{V}_{ggq}(S_2, \mathcal{J}_g(4, 5))], \quad (8)$$

$$\mathcal{J}_d(2, 3, 4) = P_q(p_{234}) [\mathcal{V}_{qZq}(\mathcal{J}_d(2, 4), S_3) + \mathcal{V}_{qZq}(\mathcal{J}_d(2, 3), S_4)]. \quad (9)$$

The stage compiler does not expand these expressions into diagrams. It collects all interaction records whose output index is the same current and compiles the sum into the output components of that current. If two different splits produce the same current index, their contributions are accumulated inside the same compiled stage output.

The four-label stage then builds the endpoint current used by the closure:

$$E_4 : \quad \mathcal{J}_d(2, 3, 4, 5) = P_q(p_{2345}) \sum_{\text{compatible splits}} \mathcal{V}(\mathcal{J}_{\text{left}}, \mathcal{J}_{\text{right}}). \quad (10)$$

The sum includes, for example, gluon-first and Z -insertion contributions such as

$$\mathcal{V}_{qZq}(\mathcal{J}_d(2, 4, 5), S_3), \quad \mathcal{V}_{qgq}(\mathcal{J}_d(2, 3, 4), S_5), \quad \mathcal{V}_{qgq}(\mathcal{J}_d(2, 3), \mathcal{J}_g(4, 5)).$$

Those are still local vertex applications selected by the model and color engine. The result is one stored current row, identified by the full CurrentIndex and by its value-slot range in the artifact.

Finally, a compiled amplitude evaluator returns raw complex roots:

$$E_{\text{amp}} : \quad R_a = \mathcal{V}_{\text{close}}(\mathcal{J}_d(2, 3, 4, 5), S_1), \quad (11)$$

with one output for each retained root. Roots with the same $(H, \text{color sector}, O)$ coherent-group key are summed before squaring. The reported leading-color matrix element is

$$|\mathcal{M}|_{\text{LC}}^2 = \mathcal{N}_{\text{AC}} \sum_{\text{groups } g} w_g \left| \sum_{a \in g} R_a \right|^2. \quad (12)$$

The same example as a DAG

A stripped-down adjacency-list representation of this example is:

```
sources interpreted by Rusticol:
S1 = J(pid=-1, labels={1}, H=b1, color=qbar-line)
S2 = J(pid= 1, labels={2}, H=b2, color=q-line)
S3 = J(pid=23, labels={3}, H=b3, color=singlet)
S4 = J(pid=21, labels={4}, H=b4, color=line gluon)
S5 = J(pid=21, labels={5}, H=b5, color=line gluon)

compiled stage E2, subset_size=2:
I0: S4 + S5 -- V3g, Pg --> J(pid=21, labels={4,5})
I1: S4 + S5 -- VggT --> J(aux=T, labels={4,5})
I2: S2 + S4 -- Vqgq,Pq --> J(pid=1, labels={2,4})
I3: S2 + S3 -- VqZq,Pq --> J(pid=1, labels={2,3})

compiled stage E3, subset_size=3:
I4: J(2,4) + S5 -- Vqgq,Pq --> J(pid=1, labels={2,4,5})
I5: S2 + J(4,5) -- Vqgq,Pq --> J(pid=1, labels={2,4,5})
I6: J(2,4) + S3 -- VqZq,Pq --> J(pid=1, labels={2,3,4})
I7: J(2,3) + S4 -- Vqgq,Pq --> J(pid=1, labels={2,3,4})

compiled stage E4, subset_size=4:
I8: J(2,4,5) + S3 -- VqZq,Pq --> J(pid=1, labels={2,3,4,5})
I9: J(2,3,4) + S5 -- Vqgq,Pq --> J(pid=1, labels={2,3,4,5})
I10: J(2,3) + J(4,5) -- Vqgq,Pq --> J(pid=1, labels={2,3,4,5})

compiled amplitude stage:
A0: J(pid=1, labels={2,3,4,5}) + S1 -- close --> R0
```

This listing is schematic: the real artifact contains all retained helicity ancestries, both leading-color sectors, source dimensions, output component ranges, momentum-slot ids, couplings, and color weights. The important point is the division of labor. The evaluator E_s performs the tensor contractions and the algebraic sum over interactions in stage s . Rusticol interprets the DAG schedule: fill sources, compute momentum sums, call E_2 , scatter its outputs into current storage, call E_3 , scatter again, and so on until E_{amp} and the final coherent square-sum.

5 Rusticol: the Implemented Rust Runtime

The previous staged PYAMPLICOL runtime proved that the shared-current DAG could match Fortran AMPLICOL numerically, but its Python loop still moved current tables between many evaluator calls. The current implementation keeps Python as the process generator and validation driver, but moves the hot eager-DAG sweep into a PyO3 extension named rusticol. The generated artifact is now a self-contained process directory used by both runtimes:

```

rt_py = pyamplicol.load_process(process_dir, runtime="python")
values_py = rt_py.evaluate(momenta)

import rusticol
rt = rusticol.Runtime.load(process_dir)
values = rt.evaluate(momenta)
profile = rt.profile(momenta)

```

The Python and Rust runtimes read the same `process_manifest.json`, the same Symbolica evaluator manifests, and the same validation momenta. The Python path is kept as a transparent reference implementation. The Rust path is the performance path.

A process directory contains four classes of data. First, it records process and model metadata: external PDG order, momentum conventions, particle and vertex tables, couplings, masses, symmetry factors, and leading-color normalization. Second, it stores the shared-current DAG: source currents, current ids, dimensions, interaction stages, output slots, retained amplitude records, helicity weights, and color factors. Third, it stores the compiled Symbolica evaluator artifacts for each stage and for the final amplitude/raw-sum stage. Fourth, it contains standalone validation assets: `validation_momenta.json` with decimal-string momenta and a `check_standalone.py` script that imports only `rusticol`, plus `numpy` when available for the optimized f64 input path.

```

# Generated by pyAmpliCol.
process_dir/
  process_manifest.json
  manifest.json
  validation_momenta.json
  check_standalone.py
  compiled/           # C++/ASM shared libraries when present
  ... evaluator states ...

```

`rusticol.Runtime.load` validates the schema, checks that the current process family is supported, loads every Symbolica evaluator, computes direct state offsets for stage outputs, allocates reusable scratch buffers, and builds a fixed execution plan. The optimized f64 path accepts NumPy arrays with shape `(batch, n_external, 4)`. For each batch, Rust fills source wavefunctions and momentum sums, executes stage evaluators in topological order, writes the returned current components into the contiguous state table, executes the amplitude/raw-sum evaluator, applies the Fortran-matched normalization, and returns a NumPy array of matrix elements.

```

for point in batch {
    fill_source_currents(state_row, point);
    fill_momentum_sums(state_row, point);
}

for stage in stages {
    evaluated = stage.evaluator.evaluate_batch(state);
    assign_stage_outputs(state, evaluated, stage.output_offsets);
}

raw = amplitude_stage.evaluate_raw_sums(state);
me2 = normalization * raw;

```

The `profile` API reports the same split used in the benchmarks: source fill, momentum setup, stage evaluator time, output assignment, amplitude evaluator time, final reduction, total wall time, and memory snapshots where the platform exposes them.

The precision API follows Symbolica’s model. `evaluate` and `evaluate_with_prec(..., 16)` use the specialized native f64 path. `evaluate_with_prec(..., 32)` maps the serialized exact evaluator state to Symbolica’s `DoubleFloat` coefficients. Other positive precisions map to arbitrary-precision `Float` evaluators. The high-precision routes share the same stage orchestration code through generic Rust traits, while the f64 route remains specialized and avoids high-precision allocation overhead.

The artifact schema is process-generic. The staged execution format no longer needs a Rust-side process-family branch: each process directory supplies the source layout, current slots, stage records, amplitude roots, coherent groups, and evaluator artifacts needed by `Rusticol`. If a model vertex or color rule is not implemented yet, generation fails with an explicit unsupported-kernel diagnostic rather than falling back to a hidden family path.

6 Generic LC Performance Matrix

Each cell compares generated-library AMPLICOL against PYAMPLICOL staged-DAG JIT and staged-DAG C++ O3 for the same final-state multiplicity n . In each cell, the upper line is generation time; the lower line is wall runtime per phase-space point. The left slot is the generated-library AMPLICOL reference value, the middle slot is PYAMPLICOL JIT, and the right slot is PYAMPLICOL C++ O3. Runtime multipliers are shown as (wall|core). When no AMPLICOL reference exists but PYAMPLICOL ran, the PYAMPLICOL slots show absolute generation time and parenthesized wall|core runtime in microseconds. Colored ratios are PYAMPLICOL/AMPLICOL: green below one, orange below two, red otherwise. When rows are generated with -validate, PYAMPLICOL JIT and C++ O3 values are compared to the same-point AMPLICOL value with a relative tolerance of 10^{-8} .

A red **VALIDATION FAILED** marker is shown next to any validated entry whose relative difference exceeds 10^{-8} .

ID base process	n=1	n=2	n=3
1 $d\bar{d} \rightarrow Z + (n-1)g$	2.904 s (x0.092) (x0.334) 0.0224 us (x6.10 x1.03) (x6.39 x1.27)	2.722 s (x0.103) (x0.556) 0.127 us (x2.68 x1.37) (x2.81 x1.50)	2.539 s (x0.192) (x1.09) 0.458 us (x2.10 x1.48) (x1.92 x1.35)
2 $u\bar{d} \rightarrow W^+ + (n-1)g$	2.734 s (x0.083) (x0.352) 0.0134 us (x7.56 x0.885) (x7.63 x1.02)	2.551 s (x0.1) (x0.533) 0.0757 us (x2.97 x1.25) (x3.12 x1.35)	2.47 s (x0.096) (x0.823) 0.261 us (x2.14 x1.39) (x2.12 x1.38)
3 $d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A	2.555 s (x0.088) (x0.557) 0.0699 us (x3.07 x1.17) (x2.69 x1.05)	2.468 s (x0.087) (x0.727) 0.179 us (x2.09 x1.19) (x2.15 x1.25)
4 $u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A	2.507 s (x0.086) (x0.622) 0.0327 us (x3.83 x0.822) (x3.58 x0.766)	2.368 s (x0.083) (x0.629) 0.0624 us (x3.05 x1.04) (x3.12 x1.13)
5 $d\bar{d} \rightarrow ZZ + (n-2)g$	N/A	2.694 s (x0.108) (x0.619) 0.166 us (x2.73 x1.41) (x2.80 x1.53)	2.5 s (x0.135) (x1.31) 0.746 us (x1.74 x1.29) (x1.84 x1.34)
6 $gg \rightarrow t\bar{t} + (n-2)g$	N/A	2.693 s (x0.103) (x0.579) 0.125 us (x3.50 x1.76) (x3.49 x1.89)	3.07 s (x0.102) (x1.10) 0.714 us (x1.90 x1.42) (x1.89 x1.43)
7 $d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A	2.616 s (x0.098) (x0.493) 0.0539 us (x4.06 x1.08) (x4.32 x1.32)	2.604 s (x0.099) (x0.836) 0.254 us (x2.43 x1.49) (x2.50 x1.58)
8 $gg \rightarrow gg + (n-2)g$	N/A	2.976 s (x0.095) (x0.562) 0.185 us (x1.85 x1.09) (x1.94 x1.18)	2.828 s (x0.13) (x1.72) 0.937 us (x1.81 x1.53) (x1.82 x1.55)
9 $d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A	N/A	2.97 s (x0.182) (x2.15) 1.59 us (x1.60 x1.25) (x1.53 x1.17)
10 $d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A	N/A	N/A
11 $d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A	N/A	N/A
12 $d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A	N/A	N/A
13 $d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A	N/A	N/A
14 $d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	N/A	N/A

Generic LC Performance Matrix (continued)

ID	base process	n=4			n=5			n=6		
∞	1 $d\bar{d} \rightarrow Z + (n-1)g$	2.988 s	(x0.165)	(x1.73)	2.559 s	(x0.221)	(x6.44)	3.577 s	(x0.283)	(x13.4)
		1.525 us	(x1.76 x1.44)	(x1.54 x1.27)	4.456 us	(x2.49 x2.27)	(x1.48 x1.27)	13.3 us	(x3.83 x3.71)	(x1.52 x1.39)
	2 $u\bar{d} \rightarrow W^+ + (n-1)g$	2.762 s	(x0.136)	(x1.33)	2.359 s	(x0.175)	(x4.32)	2.926 s	(x0.248)	(x11.8)
		1.029 us	(x1.72 x1.37)	(x1.54 x1.24)	2.881 us	(x2.45 x2.25)	(x1.59 x1.39)	9.365 us	(x3.67 x3.51)	(x1.64 x1.48)
	3 $d\bar{d} \rightarrow e^+e^- + (n-2)g$	2.74 s	(x0.123)	(x0.823)	2.352 s	(x0.132)	(x2.12)	2.562 s	(x0.148)	(x6.08)
		0.492 us	(x1.86 x1.34)	(x1.62 x1.19)	1.369 us	(x2.20 x1.91)	(x1.64 x1.37)	4.011 us	(x2.36 x2.18)	(x1.55 x1.37)
	4 $u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	2.709 s	(x0.093)	(x0.617)	2.254 s	(x0.042)	(x1.35)	2.398 s	(x0.081)	(x3.80)
		0.169 us	(x2.47 x1.44)	(x2.28 x1.37)	0.621 us	(x2.59 x2.19)	(x1.92 x1.53)	2.133 us	(x2.13 x1.93)	(x1.79 x1.58)
	5 $d\bar{d} \rightarrow ZZ + (n-2)g$	2.955 s	(x0.231)	(x2.30)	2.765 s	(x0.27)	(x6.84)	3.701 s	(x0.4)	(x22.5)
		2.271 us	(x1.66 x1.36)	(x1.35 x1.11)	6.385 us	(x5.16 x4.97)	(x1.18 x1.02)	17.69 us	(x4.75 x4.62)	(x2.92 x2.80)
	6 $gg \rightarrow t\bar{t} + (n-2)g$	4.446 s	(x0.129)	(x2.32)	7.772 s	(x0.091)	(x4.30)	22.56 s	(x0.084)	(x4.84)
		2.538 us	(x1.88 x1.56)	(x1.72 x1.46)	7.971 us	(x2.37 x2.21)	(x1.77 x1.60)	22.74 us	(x2.75 x2.61)	(x1.76 x1.65)
	7 $d\bar{d} \rightarrow t\bar{t} + (n-2)g$	3.216 s	(x0.123)	(x1.52)	3.161 s	(x0.137)	(x5.26)	4.852 s	(x0.176)	(x6.20)
		0.93 us	(x2.38 x1.90)	(x2.18 x1.76)	2.901 us	(x3.10 x2.81)	(x2.10 x1.81)	9.349 us	(x7.07 x6.89)	(x1.94 x1.83)
	8 $gg \rightarrow gg + (n-2)g$	2.915 s	(x0.293)	(x4.95)	3.092 s	(x0.276)	(x22.6)	5.214 s	(x0.397)	(x39.8)
		3.468 us	(x2.13 x1.88)	(x1.77 x1.58)	10.6 us	(x3.20 x3.07)	(x4.11 x3.95)	31.12 us	(x3.00 x2.91)	(x3.67 x3.58)
∞	9 $d\bar{d} \rightarrow ZZZ + (n-3)g$	4.394 s	(x0.5)	(x3.61)	5.062 s	(x0.233)	(x12.2)	6.539 s	(x0.451)	(x26.7)
		5.447 us	(x1.70 x1.41)	(x1.10 x0.911)	13.21 us	(x5.26 x5.11)	(x3.03 x2.90)	34.28 us	(x7.11 x6.90)	(x3.01 x2.93)
	10 $d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	5.506 s	(x0.097)	(x0.657)	5.482 s	(x0.038)	(x1.46)	5.287 s	(x0.082)	(x3.29)
		0.8 us	(x2.49 x1.93)	(x2.03 x1.62)	1.619 us	(x2.50 x2.22)	(x1.93 x1.64)	4.002 us	(x2.27 x2.08)	(x1.55 x1.37)
	11 $d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	4.068 s	(x0.103)	(x1.23)	6.678 s	(x0.083)	(x3.22)	9.256 s	(x0.109)	(x8.82)
		1.138 us	(x1.93 x1.55)	(x1.69 x1.34)	3.474 us	(x6.76 x6.49)	(x1.89 x1.63)	9.971 us	(x9.42 x9.24)	(x5.21 x5.04)
∞	12 $d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	5.089 s	(x0.103)	(x0.762)	5.582 s	(x0.036)	(x1.42)	5.336 s	(x0.081)	(x3.09)
		1.634 us	(x1.11 x0.901)	(x0.941 x0.765)	2.845 us	(x1.37 x1.23)	(x1.08 x0.931)	6.869 us	(x1.79 x1.70)	(x0.849 x0.761)
∞	13 $d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	2.858 s	(x0.084)	(x0.624)	3.468 s	(x0.077)	(x1.13)	6.356 s	(x0.088)	(x1.54)
		0.187 us	(x2.01 x1.08)	(x1.96 x1.10)	0.69 us	(x2.13 x1.74)	(x1.76 x1.38)	2.334 us	(x2.41 x2.19)	(x1.61 x1.43)
∞	14 $d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A			N/A			N/A	2.869 s	35.59 s
								N/A	(58 54)	(14 10.9)

Generic LC Performance Matrix (continued)

ID	base process	n=7			n=8			n=9		
6	1 $d\bar{d} \rightarrow Z + (n-1)g$	4.506 s	(x0.57)	(x33.2)	9.224 s	(x0.946)	(N/A)	24.26 s	(x1.51)	(N/A)
		36.49 us	(x3.88 x3.81)	(x1.58 x1.49)	91.02 us	(x4.29 x4.24)	(N/A N/A)	222 us	(x6.36 x6.30)	(N/A N/A)
	2 $u\bar{d} \rightarrow W^+ + (n-1)g$	3.157 s	(x0.587)	(x34.4)	4.631 s	(x1.38)	(N/A)	9.949 s	(x3.10)	(N/A)
		26.94 us	(x3.12 x3.02)	(x1.59 x1.48)	68.3 us	(x3.74 x3.67)	(N/A N/A)	167 us	(x5.27 x5.21)	(N/A N/A)
	3 $d\bar{d} \rightarrow e^+e^- + (n-2)g$	2.826 s	(x0.292)	(x15.5)	3.74 s	(x0.639)	(N/A)	7.283 s	(x1.08)	(N/A)
		11.84 us	(x2.58 x2.44)	(x1.54 x1.40)	32.37 us	(x3.39 x3.30)	(N/A N/A)	83.59 us	(x3.97 x3.91)	(N/A N/A)
	4 $u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	2.409 s	(x0.2)	(x11.3)	2.656 s	(x0.531)	(N/A)	3.506 s	(x1.53)	(N/A)
		7.296 us	(x2.47 x2.34)	(x1.62 x1.49)	20.77 us	(x3.48 x3.38)	(N/A N/A)	55.81 us	(x3.57 x3.52)	(N/A N/A)
	5 $d\bar{d} \rightarrow ZZ + (n-2)g$	6.599 s	(x0.534)	(x21.3)	16.65 s	(x0.573)	(N/A)	27.71 s	(x0.968)	(N/A)
		44.03 us	(x4.35 x4.29)	(x1.42 x1.32)	105 us	(x6.57 x6.49)	(N/A N/A)	250 us	(x7.38 x7.30)	(N/A N/A)
	6 $gg \rightarrow t\bar{t} + (n-2)g$	182 s	(x0.029)	(x1.71)	2.7e+03 s	(x0.004)	(N/A)	679 s	(x0.08)	(N/A)
		61.03 us	(x4.38 x4.31)	(x2.54 x2.48)	178 us	(x2.75 x2.71)	(N/A N/A)	465 us	(x7.39 x7.36)	(N/A N/A)
	7 $d\bar{d} \rightarrow t\bar{t} + (n-2)g$	9.569 s	(x0.28)	(x17.0)	75.77 s	(x0.356)	(N/A)	768 s	(x0.04)	(N/A)
		26.07 us	(x5.65 x5.54)	(x3.49 x3.37)	66.74 us	(x8.40 x8.33)	(N/A N/A)	177 us	(x3.56 x3.52)	(N/A N/A)
	8 $gg \rightarrow gg + (n-2)g$	8.765 s	(x2.40)	(x48.2)	83.6 s	(x0.532)	(N/A)	583 s	(N/A)	(N/A)
		86.06 us	(x6.03 x5.96)	(x2.33 x2.28)	240 us	(x8.30 x8.23)	(N/A N/A)	512 us	(N/A N/A)	(N/A N/A)
7	9 $d\bar{d} \rightarrow ZZZ + (n-3)g$	15.02 s	(x0.491)	(N/A)	75.75 s	(x0.273)	(N/A)	152 s	(x0.372)	(N/A)
		82.46 us	(x5.34 x5.27)	(N/A N/A)	191 us	(x13.4 x13.2)	(N/A N/A)	456 us	(x16.7 x16.7)	(N/A N/A)
	10 $d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	6.741 s	(x0.139)	(x6.38)	17.32 s	(x0.144)	(N/A)	34.33 s	(x0.166)	(N/A)
		10.06 us	(x4.20 x4.07)	(x2.25 x2.12)	24.71 us	(x5.42 x5.29)	(N/A N/A)	60.7 us	(x3.97 x3.91)	(N/A N/A)
	11 $d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	25.46 s	(x0.112)	(N/A)	180 s	(x0.044)	(N/A)	1.57e+03 s	(x0.017)	(N/A)
		28.33 us	(x9.50 x9.40)	(N/A N/A)	74.86 us	(x10.1 x9.99)	(N/A N/A)	225 us	(x13.4 x13.3)	(N/A N/A)
	12 $d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	7.641 s	(x0.114)	(x5.13)	19 s	(x0.115)	(N/A)	41.03 s	(x0.125)	(N/A)
		15.49 us	(x2.36 x2.28)	(x1.84 x1.76)	37.34 us	(x2.88 x2.81)	(N/A N/A)	84.88 us	(x2.38 x2.34)	(N/A N/A)
	13 $d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	14.71 s	(x0.05)	(x2.30)	41.61 s	(x0.051)	(N/A)	156 s	(x0.058)	(N/A)
		7.633 us	(x4.94 x4.79)	(x2.96 x2.81)	24.73 us	(x6.27 x6.16)	(N/A N/A)	60.86 us	(x5.37 x5.30)	(N/A N/A)
	14 $d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A	25.5 s	N/A	N/A			N/A		
		N/A	(700 682)	N/A						

Matrix run settings. The matrix table is generated with one uniform benchmark setup. The full refresh command is:

```
pyAmpliCol/dependencies/.venv/bin/python pyAmpliCol/docs/result_matrix.py --color-accuracy lc --data
/Users/vjhirsch/HEP_programs/AmpliCol/pyAmpliCol/docs/result_matrix_data.json --table
/Users/vjhirsch/HEP_programs/AmpliCol/pyAmpliCol/docs/result_matrix_table.tex --generate-data 1-9 --show-data
1-9 --time-limit 900 --target-runtime 1 --amplicol-points 10000 --jobs 4 --n-cores 4 --process-workers 1
--columns-per-table 3
```

Partial refreshes use the same settings with `-base-process`, `-skip-jit`, `-skip-cpp-o3`, `-process-ids`, or `-skip-amplicol` added only to restrict which cells are updated. PyAmpliCol-only refreshes may additionally use `-process-workers 5`; this is intentionally disabled when the AMPLICOL reference is refreshed because that path uses shared top-level generated-library files. AMPLICOL reference timings use the generated-library sequence `amplicol_generate -library=create, make -j4 amplicol_generate_library`, and a direct generated-library benchmark using `-library=use`; they do not use integration-driver timings. PYAMPLICOL rows use the generic LC staged-DAG artifact, Rusticol double precision, JIT or C++ O3 Symbolica evaluators as labelled, and `time-process` with a one-second target runtime. The `-time-limit` setting applies only to C++ O3 generation; AMPLICOL and JIT rows are left uncapped.

Structural reference limitations. For the four-quark-line row $d\bar{d} \rightarrow u\bar{u} s\bar{s} c\bar{c} + (n-6)g$, the AMPLICOL reference column is a structural N/A: Fortran AMPLICOL stops with **Unknown number of quarks and anti-quarks**. The PYAMPLICOL cells are therefore shown as absolute generation/runtime timings when they are available.

C++ O3 compile-budget gaps. The following cells have validated AMPLICOL and PYAMPLICOL JIT rows but no completed C++ O3 row; their C++ slot is a localized performance-generation compile-budget gap, not a physics-validation failure: $d\bar{d} \rightarrow Z + (n-1)g$, $n=8$: not run; $d\bar{d} \rightarrow Z + (n-1)g$, $n=9$: not run; $u\bar{d} \rightarrow W^+ + (n-1)g$, $n=8$: not run; $u\bar{d} \rightarrow W^+ + (n-1)g$, $n=9$: not run; $d\bar{d} \rightarrow e^+e^- + (n-2)g$, $n=8$: not run; $d\bar{d} \rightarrow e^+e^- + (n-2)g$, $n=9$: not run; $u\bar{d} \rightarrow e^+\nu_e + (n-2)g$, $n=8$: not run; $u\bar{d} \rightarrow e^+\nu_e + (n-2)g$, $n=9$: not run; plus 19 more localized entries.

Documented JIT backend limitations. The following PYAMPLICOL JIT entries are rendered as N/A because they hit a localized Symbolica/SymJIT AArch64 complex-JIT materialization limitation, not a generated-library comparison or physics-validation failure: $gg \rightarrow gg + (n-2)g$, $n=9$, PYAMPLICOL JIT; $d\bar{d} \rightarrow u\bar{u} s\bar{s} c\bar{c} + (n-6)g$, $n=8$, PYAMPLICOL JIT; $d\bar{d} \rightarrow u\bar{u} s\bar{s} c\bar{c} + (n-6)g$, $n=9$, PYAMPLICOL JIT. The raw assertion is retained in `result_matrix_data.json`. The current managed dependency set uses SymJIT 2.19.3, which fixes the smaller historical v220 AArch64 JIT segfault reproduced by `MRE_symjit_bug_new.py` and the pure-gluon $n=4$ NLC/full JIT materialization case. Cached LC backend-limit entries should be refreshed before being interpreted under this newer SymJIT version.

7 Generic NLC Performance Matrix

Each cell compares colour-matrix AMPLI_{COL} probes against PYAMPLI_{COL} staged-DAG JIT and staged-DAG C++ O3 for the same final-state multiplicity n . In each cell, the upper line is generation time; the lower line is wall runtime per phase-space point. The left slot is the generated-library AMPLI_{COL} reference value, the middle slot is PYAMPLI_{COL} JIT, and the right slot is PYAMPLI_{COL} C++ O3. Runtime multipliers are shown as (wall|core). When no AMPLI_{COL} reference exists but PYAMPLI_{COL} ran, the PYAMPLI_{COL} slots show absolute generation time and parenthesized wall|core runtime in microseconds. For NLC/full-colour tables, AMPLI_{COL} reference values come from the dedicated `amplicol_color_probe` imode-2 colour-matrix driver rather than the LC generated-library timing driver. Colored ratios are PYAMPLI_{COL}/AMPLI_{COL}: green below one, orange below two, red otherwise. When rows are generated with `-validate`, PYAMPLI_{COL} JIT and C++ O3 values are compared to the same-point AMPLI_{COL} value with a relative tolerance of 10^{-8} .

A red **VALIDATION FAILED** marker is shown next to any validated entry whose relative difference exceeds 10^{-8} .

ID	base process	n=1			n=2			n=3		
1	$d\bar{d} \rightarrow Z + (n-1)g$	0.155 s	(x1.43)	(x9.49)	0.154 s	(x1.42)	(x10.4)	0.159 s	(x2.59)	(x26.9)
		11.76 us	(x0.015 x0.003)	(x0.014 x0.003)	27.55 us	(x0.014 x0.007)	(x0.013 x0.007)	65.3 us	(x0.028 x0.02)	(x0.027 x0.02)
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	0.156 s	(x1.32)	(x5.93)	0.147 s	(x1.42)	(x9.15)	0.138 s	(x1.88)	(x19.7)
		12.18 us	(x0.01 x0.001)	(x0.011 x0.002)	26.39 us	(x0.009 x0.004)	(x0.009 x0.004)	63.77 us	(x0.017 x0.012)	(x0.018 x0.012)
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A			0.16 s	(x1.36)	(x8.31)	0.142 s	(x1.76)	(x13.0)
					17.58 us	(x0.011 x0.004)	(x0.011 x0.004)	37.56 us	(x0.01 x0.006)	(x0.011 x0.006)
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A			0.15 s	(x1.36)	(x8.57)	0.198 s	(x1.92)	(x8.09)
					17.28 us	(x0.007 x0.001)	(x0.007 x0.001)	39.79 us	(x0.006 x0.002)	(x0.005 x0.002)
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A			0.146 s	(x1.54)	(x10.8)	0.144 s	(x2.13)	(x23.6)
					40.03 us	(x0.011 x0.006)	(x0.011 x0.006)	88.3 us	(x0.015 x0.011)	(x0.014 x0.011)
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A			0.147 s	(x1.60)	(x14.1)	0.141 s	(x4.41)	(x136)
					18.5 us	(x0.047 x0.028)	(x0.048 x0.029)	52.62 us	(x0.37 x0.318)	(x0.352 x0.303)
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A			0.147 s	(x1.57)	(x9.91)	0.142 s	(x2.35)	(x37.0)
					18.2 us	(x0.025 x0.01)	(x0.026 x0.011)	45.51 us	(x0.06 x0.04)	(x0.058 x0.038)
8	$gg \rightarrow gg + (n-2)g$	N/A			0.151 s	(x1.94)	(x29.2)	0.149 s	(x11.5)	(x578)
					20.41 us	(x0.098 x0.075)	(x0.101 x0.074)	92.54 us	(x0.492 x0.45)	(x0.466 x0.425)
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A			N/A			0.159 s	(x2.29)	(x36.6)
								149 us	(x0.017 x0.014)	(x0.016 x0.012)
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A			N/A			N/A		
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A			N/A			N/A		
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A			N/A			N/A		
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A			N/A			N/A		
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A			N/A			N/A		

Generic NLC Performance Matrix (continued)

ID	base process	n=4		
1	$d\bar{d} \rightarrow Z + (n-1)g$	0.151 s	(x8.23)	(x329)
		221 us	(x0.208 x0.189)	(x0.204 x0.188)
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	0.154 s	(x5.13)	(x146)
		222 us	(x0.08 x0.067)	(x0.075 x0.064)
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	0.147 s	(x2.09)	(x28.4)
		90.24 us	(x0.019 x0.015)	(x0.02 x0.015)
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	0.146 s	(x1.66)	(x17.1)
		89.27 us	(x0.008 x0.005)	(x0.008 x0.006)
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	0.148 s	(x4.53)	(x121)
		250 us	(x0.032 x0.026)	(x0.03 x0.023)
6	$gg \rightarrow t\bar{t} + (n-2)g$	0.145 s	(x56.6)	timeout
		332 us	(x1.22 x1.16)	(N/A N/A)
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	0.151 s	(x11.7)	(x841)
		179 us	(x0.432 x0.398)	(x0.421 x0.39)
8	$gg \rightarrow gg + (n-2)g$	0.151 s	(x299)	(N/A)
		1.07e+03 us	(x0.853 x0.783)	(N/A N/A)
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	0.158 s	(x4.32)	(x114)
		368 us	(x0.024 x0.02)	(x0.019 x0.015)
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	0.156 s	(x2.03)	(x29.3)
		77.65 us	(x0.023 x0.018)	(x0.022 x0.018)
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	0.152 s	(x3.42)	(x70.7)
		87.79 us	(x0.058 x0.047)	(x0.051 x0.039)
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	0.167 s	(x1.95)	(x29.2)
		130 us	(x0.014 x0.012)	(x0.013 x0.011)
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A		
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A		

Matrix run settings. The matrix table is generated with one uniform benchmark setup. The full refresh command is:

```
pyAmpliCol/dependencies/.venv/bin/python pyAmpliCol/docs/result_matrix.py --color-accuracy nlc --data
/Users/vjhirsch/HEP_programs/AmpliCol/pyAmpliCol/docs/result_matrix_nlc_data.json --table
/Users/vjhirsch/HEP_programs/AmpliCol/pyAmpliCol/docs/result_matrix_nlc_table.tex --generate-data 1-9
--show-data 1-9 --time-limit 900 --target-runtime 1 --amplicol-points 10000 --jobs 4 --n-cores 4
--process-workers 1 --columns-per-table 3
```

Partial refreshes use the same settings with `-base-process`, `-skip-jit`, `-skip-cpp-o3`, `-process-ids`, or `-skip-amplicol` added only to restrict which cells are updated. PyAmpliCol-only refreshes may additionally use `-process-workers 5`; this is intentionally disabled when the AMPLICOL reference is refreshed because that path uses shared top-level generated-library files. AMPLICOL reference timings for this colour table currently use the dedicated `amplicol_color_probe` imode-2 colour-matrix driver, not integration-driver timings. The driver exposes the same absolute AMPLICOL NLC/full-colour matrix element and splits amplitude-evaluation from colour-contraction time, but it is not yet a generated-library driver. PYAMPLICOL rows use the generic staged-DAG artifact, Rusticol double precision, JIT or C++ O3 Symbolica evaluators as labelled, and `time-process` with a one-second target runtime. The `-time-limit` setting applies only to C++ O3 generation; AMPLICOL and JIT rows are left uncapped.

C++ O3 compile-budget gaps. The following cells have validated AMPLICOL and PYAMPLICOL JIT rows but no completed C++ O3 row; their C++ slot is a localized performance-generation compile-budget gap, not a physics-validation failure: $gg \rightarrow gg + (n-2)g$, $n = 4$: not run; $gg \rightarrow t\bar{t} + (n-2)g$, $n = 4$: timed out.

8 Generic Full-Colour Performance Matrix

Each cell compares colour-matrix AMPLI_{COL} probes against PYAMPLI_{COL} staged-DAG JIT and staged-DAG C++ O3 for the same final-state multiplicity n . In each cell, the upper line is generation time; the lower line is wall runtime per phase-space point. The left slot is the generated-library AMPLI_{COL} reference value, the middle slot is PYAMPLI_{COL} JIT, and the right slot is PYAMPLI_{COL} C++ O3. Runtime multipliers are shown as (wall|core). When no AMPLI_{COL} reference exists but PYAMPLI_{COL} ran, the PYAMPLI_{COL} slots show absolute generation time and parenthesized wall|core runtime in microseconds. For NLC/full-colour tables, AMPLI_{COL} reference values come from the dedicated `amplicol_color_probe` imode-2 colour-matrix driver rather than the LC generated-library timing driver. Colored ratios are PYAMPLI_{COL}/AMPLI_{COL}: green below one, orange below two, red otherwise. When rows are generated with `-validate`, PYAMPLI_{COL} JIT and C++ O3 values are compared to the same-point AMPLI_{COL} value with a relative tolerance of 10^{-8} .

A red **VALIDATION FAILED** marker is shown next to any validated entry whose relative difference exceeds 10^{-8} .

ID	base process	n=1			n=2			n=3		
1	$d\bar{d} \rightarrow Z + (n-1)g$	0.182 s	(x1.53)	(x6.58)	0.149 s	(x2.70)	(x10.4)	0.157 s	(x2.41)	(x26.0)
		12.25 us	(x0.014 x0.003)	(x0.014 x0.004)	26.39 us	(x0.013 x0.007)	(x0.014 x0.007)	66.05 us	(x0.027 x0.02)	(x0.028 x0.021)
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	0.174 s	(x1.60)	(x6.19)	0.14 s	(x1.52)	(x9.59)	0.152 s	(x1.72)	(x20.5)
		12.62 us	(x0.017 x0.002)	(x0.011 x0.002)	26.8 us	(x0.009 x0.004)	(x0.01 x0.004)	64.58 us	(x0.017 x0.012)	(x0.017 x0.011)
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	N/A			0.169 s	(x1.33)	(x7.77)	0.152 s	(x1.46)	(x12.4)
					18.21 us	(x0.01 x0.004)	(x0.011 x0.004)	38.06 us	(x0.01 x0.006)	(x0.011 x0.006)
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	N/A			0.146 s	(x1.39)	(x8.44)	0.145 s	(x1.42)	(x11.3)
					16.93 us	(x0.007 x0.001)	(x0.007 x0.001)	36.09 us	(x0.005 x0.002)	(x0.006 x0.002)
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	N/A			0.145 s	(x1.52)	(x10.8)	0.152 s	(x1.90)	(x21.4)
					38.12 us	(x0.011 x0.006)	(x0.012 x0.006)	88.42 us	(x0.015 x0.011)	(x0.014 x0.011)
6	$gg \rightarrow t\bar{t} + (n-2)g$	N/A			0.144 s	(x1.82)	(x14.1)	0.14 s	(x4.24)	(x137)
					19.01 us	(x0.047 x0.028)	(x0.047 x0.028)	54.94 us	(x0.342 x0.292)	(x0.34 x0.29)
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	N/A			0.143 s	(x1.72)	(x10.0)	0.14 s	(x2.35)	(x35.0)
					18.06 us	(x0.025 x0.01)	(x0.026 x0.011)	45.45 us	(x0.059 x0.039)	(x0.056 x0.037)
8	$gg \rightarrow gg + (n-2)g$	N/A			0.148 s	(x1.85)	(x29.4)	0.131 s	(x14.8)	(x645)
					21.71 us	(x0.086 x0.066)	(x0.089 x0.068)	106 us	(x0.431 x0.379)	(x0.43 x0.349)
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	N/A			N/A			0.112 s	(x2.90)	(x49.9)
								142 us	(x0.017 x0.014)	(x0.018 x0.013)
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	N/A			N/A			N/A		
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	N/A			N/A			N/A		
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	N/A			N/A			N/A		
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A			N/A			N/A		
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A			N/A			N/A		

Generic Full-Colour Performance Matrix (continued)

ID	base process	n=4		
1	$d\bar{d} \rightarrow Z + (n-1)g$	0.172 s 228 us	(x7.27) (x0.206 x0.184)	(x310) (x0.207 x0.188)
2	$u\bar{d} \rightarrow W^+ + (n-1)g$	0.167 s 235 us	(x4.88) (x0.077 x0.065)	(x143) (x0.081 x0.068)
3	$d\bar{d} \rightarrow e^+e^- + (n-2)g$	0.165 s 94.45 us	(x1.95) (x0.019 x0.014)	(x26.6) (x0.019 x0.015)
4	$u\bar{d} \rightarrow e^+\nu_e + (n-2)g$	0.175 s 91.49 us	(x1.61) (x0.009 x0.006)	(x14.6) (x0.008 x0.005)
5	$d\bar{d} \rightarrow ZZ + (n-2)g$	0.16 s 254 us	(x4.22) (x0.032 x0.025)	(x111) (x0.031 x0.024)
6	$gg \rightarrow t\bar{t} + (n-2)g$	0.147 s 379 us	(x56.5) (x1.09 x1.00)	(N/A) (N/A N/A)
7	$d\bar{d} \rightarrow t\bar{t} + (n-2)g$	0.168 s 188 us	(x10.6) (x0.412 x0.37)	(x803) (x0.419 x0.385)
8	$gg \rightarrow gg + (n-2)g$	0.149 s 1.71e+03 us	(x389) (x0.594 x0.468)	(N/A) (N/A N/A)
9	$d\bar{d} \rightarrow ZZZ + (n-3)g$	0.153 s 365 us	(x4.48) (x0.022 x0.018)	(x119) (x0.021 x0.017)
10	$d\bar{d} \rightarrow e^+e^-ZH + (n-4)g$	0.149 s 78.75 us	(x2.07) (x0.024 x0.019)	(x31.1) (x0.022 x0.017)
11	$d\bar{d} \rightarrow t\bar{t}ZH + (n-4)g$	0.147 s 84.01 us	(x2.88) (x0.057 x0.046)	(x69.8) (x0.052 x0.041)
12	$d\bar{d} \rightarrow e^+e^-e^+e^- + (n-4)g$	0.16 s 120 us	(x2.02) (x0.015 x0.012)	(x30.7) (x0.014 x0.012)
13	$d\bar{d} \rightarrow u\bar{u}s\bar{s} + (n-4)g$	N/A		
14	$d\bar{d} \rightarrow u\bar{u}s\bar{s}c\bar{c} + (n-6)g$	N/A		

Matrix run settings. The matrix table is generated with one uniform benchmark setup. The full refresh command is:

```
pyAmpliCol/dependencies/.venv/bin/python pyAmpliCol/docs/result_matrix.py --color-accuracy full --data
/Users/vjhirsch/HEP_programs/AmpliCol/pyAmpliCol/docs/result_matrix_full_data.json --table
/Users/vjhirsch/HEP_programs/AmpliCol/pyAmpliCol/docs/result_matrix_full_table.tex --generate-data 1-9
--show-data 1-9 --time-limit 900 --target-runtime 1 --amplicol-points 10000 --jobs 4 --n-cores 4
--process-workers 1 --columns-per-table 3
```

Partial refreshes use the same settings with `-base-process`, `-skip-jit`, `-skip-cpp-o3`, `-process-ids`, or `-skip-amplicol` added only to restrict which cells are updated. PyAmpliCol-only refreshes may additionally use `-process-workers 5`; this is intentionally disabled when the AMPLICOL reference is refreshed because that path uses shared top-level generated-library files. AMPLICOL reference timings for this colour table currently use the dedicated `amplicol_color_probe` imode-2 colour-matrix driver, not integration-driver timings. The driver exposes the same absolute AMPLICOL NLC/full-colour matrix element and splits amplitude-evaluation from colour-contraction time, but it is not yet a generated-library driver. PYAMPLICOL rows use the generic staged-DAG artifact, Rusticol double precision, JIT or C++ O3 Symbolica evaluators as labelled, and `time-process` with a one-second target runtime. The `-time-limit` setting applies only to C++ O3 generation; AMPLICOL and JIT rows are left uncapped.

C++ O3 compile-budget gaps. The following cells have validated AMPLICOL and PYAMPLICOL JIT rows but no completed C++ O3 row; their C++ slot is a localized performance-generation compile-budget gap, not a physics-validation failure: $gg \rightarrow t\bar{t} + (n-2)g$, $n = 4$: not run; $gg \rightarrow gg + (n-2)g$, $n = 4$: not run.

Archived low-multiplicity validation fixture. The $n \leq 4$ entries of the LC, NLC and full-colour process matrices are also stored as a compact integration-test fixture in `tests/fixtures/result_matrix_references.json`. The fixture contains the process string, colour accuracy, deterministic phase-space point, Fortran AMPLICOL reference value, PYAMPLICOL JIT value, tolerance and relative difference for each validated cell. It is checked by `tests/integration/test_result_matrix_references.py` without requiring a local Fortran AMPLICOL run. To refresh the fixture from the matrix caches, run `scripts/refresh_result_matrix_fixtures.py`; with `-refresh-caches`, the script first reruns the documented `result_matrix.py` validation workflow against Fortran AMPLICOL.

9 Performance Results

Tables 8 and 9 report the current benchmark state for $d\bar{d} \rightarrow Z + ng$, $n = 1, \dots, 9$. All successful `pyAmpliCol` rows have been validated against Fortran `AMPLiCol` with maximum relative differences typically of order 10^{-13} . The Fortran reference rows use the generated-library workflow `-library=create`, compiled `amplicol_generate_library`, and probe evaluation with `-library=use`. Timing values are shown in monospace font. Green, orange, and red multiplier text means faster than Fortran `AMPLiCol`, slower but below a factor two, and slower by a factor two or more. The blue rows are the Fortran `AMPLiCol` references. The green highlighted rows are the current `pyAmpliCol` production route, `rusticol` with C++ O3 staged evaluators. The dependency installer currently uses upstream Symbolica `dev` with the `pyAmpliCol` local patches and pins `SymJIT` to 2.19.3 at commit `dee77e14c8bf9c8d5304dbda2e393c52ee3e4cd4`. This fixes the historical AArch64 complex-JIT crash captured by `MRE_symjit_bug_new.py`; staged-DAG JIT rows refreshed after this dependency update are noted in the tables below.

Table 8: Staged/Rusticol performance summary for $n = 1, \dots, 4$.

n route	setup	gen [s]	wall [us/pt]	eval [us/pt]	notes
1 reference	AmpliCol	2.22	N/A	0.182	Fortran AmpliCol; with -library=use
1 Python staged	pyAmpliCol - JIT	0.016 (x0.01)	13.34(29) (x73.10)	0.607(17) (x3.33)	JIT; batch=16
1 Python staged	pyAmpliCol - C++	0.774 (x0.35)	13.01(12) (x71.29)	0.521(10) (x2.85)	C++; O3; chunk=None; batch=16
1 Python staged	pyAmpliCol - ASM	0.560 (x0.25)	14.03(30) (x76.88)	0.912(18) (x5.00)	ASM; O3; chunk=None; batch=16
1 Rusticol	pyAmpliCol - Rusticol	0.984 (x0.44)	0.412(1) (x2.26)	0.1938(2) (x1.06)	Rusticol PyO3; C++ O3 process artifact; cached NumPy input; reusable stage scratch; single-chunk direct output; batch=1024
2 reference	AmpliCol	2.17	N/A	0.520	Fortran AmpliCol; with -library=use
2 Python staged	pyAmpliCol - JIT	0.066 (x0.03)	16.15(15) (x31.04)	1.61(2) (x3.09)	JIT; batch=16
2 Python staged	pyAmpliCol - C++	1.86 (x0.86)	16.05(14) (x30.85)	1.16(2) (x2.23)	C++; O3; chunk=None; batch=16
2 Python staged	pyAmpliCol - ASM	0.940 (x0.43)	16.48(20) (x31.67)	2.16(2) (x4.15)	ASM; O3; chunk=None; batch=16
2 Rusticol	pyAmpliCol - Rusticol	2.00 (x0.92)	0.984(1) (x1.89)	0.653(1) (x1.26)	Rusticol PyO3; C++ O3 process artifact; cached NumPy input; reusable stage scratch; single-chunk direct output; batch=1024
3 reference	AmpliCol	2.29	N/A	1.61	Fortran AmpliCol; with -library=use
3 Python staged	pyAmpliCol - JIT	0.281 (x0.12)	22.67(42) (x14.08)	4.72(5) (x2.93)	JIT; batch=16
3 Python staged	pyAmpliCol - C++	5.69 (x2.48)	20.90(40) (x12.98)	3.06(4) (x1.90)	C++; O3; chunk=None; batch=16
3 Python staged	pyAmpliCol - ASM	1.52 (x0.66)	23.87(45) (x14.83)	6.17(8) (x3.83)	ASM; O3; chunk=None; batch=16
3 Rusticol	pyAmpliCol - Rusticol	5.48 (x2.39)	2.783(13) (x1.73)	2.177(10) (x1.35)	Rusticol PyO3; C++ O3 process artifact; cached NumPy input; reusable stage scratch; batch=128
4 reference	AmpliCol	2.65	N/A	5.11	Fortran AmpliCol; with -library=use
4 Python staged	pyAmpliCol - JIT	0.431 (x0.16)	14.23(1) (x2.78)	11.20(1) (x2.19)	JIT; batch=16; last successful JIT refresh
4 Python staged	pyAmpliCol - C++	20.7 (x7.80)	27.78(13) (x5.43)	7.75(6) (x1.52)	C++; O3; chunk=None; batch=16
4 Python staged	pyAmpliCol - ASM	2.04 (x0.77)	36.30(54) (x7.10)	15.80(15) (x3.09)	ASM; O3; chunk=None; batch=16
4 Rusticol	pyAmpliCol - Rusticol	19.9 (x7.50)	7.23(7) (x1.41)	6.10(6) (x1.19)	Rusticol PyO3; C++ O3 process artifact; cached NumPy input; reusable stage scratch; batch=128

Table 9: Staged/Rusticol performance summary for $n = 5, \dots, 9$.

n route	setup	gen [s]	wall [us/pt]	eval [us/pt]	notes
5 reference	AmpliCol	3.26	N/A	13.9	Fortran AmpliCol; with -library=use
5 Python staged	pyAmpliCol - JIT	0.991 (x0.30)	44.33(5) (x3.19)	40.56(4) (x2.92)	JIT; batch=16; last successful JIT refresh
5 Python staged	pyAmpliCol - C++	54.3 (x16.67)	45.06(48) (x3.25)	20.83(26) (x1.50)	C++; O3; chunk=64; batch=128; chunk compile workers=10; two-stage save/load
5 Python staged	pyAmpliCol - C++ (O2)	50.7 (x15.57)	51.04(87) (x3.68)	23.82(42) (x1.72)	C++; O2; chunk=64; batch=16
5 Python staged	pyAmpliCol - ASM	10.2 (x3.13)	75.69(88) (x5.45)	48.74(72) (x3.51)	ASM; O3; chunk=64; batch=16
5 Rusticol	pyAmpliCol - Rusticol	50.2 (x15.41)	19.63(28) (x1.41)	17.50(25) (x1.26)	Rusticol PyO3; C++ O3; chunk=64; cached NumPy input; reusable stage scratch; batch=128; two-stage save/load
6 reference	AmpliCol	4.66	N/A	37.3	Fortran AmpliCol; with -library=use
6 Python staged	pyAmpliCol - JIT	2.96 (x0.64)	163.0(5) (x4.37)	157.0(5) (x4.21)	JIT; batch=16; last successful JIT refresh
6 Python staged	pyAmpliCol - C++	151 (x32.44)	94.79(73) (x2.54)	56.36(49) (x1.51)	C++; O3; chunk=64; batch=128; chunk compile workers=10; two-stage save/load
6 Python staged	pyAmpliCol - C++ (O2)	140 (x30.07)	108.4(1.8) (x2.91)	68.5(1.6) (x1.84)	C++; O2; chunk=64; batch=16
6 Python staged	pyAmpliCol - ASM	21.4 (x4.60)	183.7(3.3) (x4.93)	141.2(2.7) (x3.79)	ASM; O3; chunk=64; batch=16
6 Rusticol	pyAmpliCol - Rusticol	141 (x30.29)	55.63(26) (x1.49)	50.12(23) (x1.34)	Rusticol PyO3; C++ O3; chunk=64; cached NumPy input; reusable stage scratch; batch=128; two-stage save/load
7 reference	AmpliCol	9.28	N/A	95.2	Fortran AmpliCol; with -library=use
7 Python staged	pyAmpliCol - JIT	8.85 (x0.95)	536(2) (x5.63)	525(2) (x5.52)	JIT; batch=16; last successful JIT refresh
7 Python staged	pyAmpliCol - C++	417 (x44.92)	259.7(1.4) (x2.73)	194.7(1.2) (x2.05)	C++; O3; chunk=96; batch=128; chunk compile workers=10; two-stage save/load
7 Python staged	pyAmpliCol - C++ (O1)	219 (x23.59)	421.6(2.4) (x4.43)	348.9(1.9) (x3.67)	C++; O1; chunk=96; batch=16
7 Python staged	pyAmpliCol - ASM	36.2 (x3.90)	384.4(2.8) (x4.04)	320.2(2.4) (x3.36)	ASM; O3; chunk=96; batch=16
7 Rusticol	pyAmpliCol - Rusticol	391 (x42.12)	197.8(4) (x2.08)	186.1(4) (x1.95)	Rusticol PyO3; C++ O3; chunk=96; cached NumPy input; reusable stage scratch; batch=128; two-stage save/load
8 reference	AmpliCol	26.1	N/A	234	Fortran AmpliCol; with -library=use
8 Python staged	pyAmpliCol - JIT	42.3 (x1.62)	1783(4) (x7.61)	1763(4) (x7.52)	JIT; batch=16; last successful JIT refresh
8 Python staged	pyAmpliCol - C++	1.06e3 (x40.69)	642.6(2.1) (x2.74)	491.8(1.4) (x2.10)	C++; O3; chunk=96; batch=128; chunk compile workers=10; two-stage save/load
8 Python staged	pyAmpliCol - C++ (O1)	577 (x22.15)	997.2(3.3) (x4.26)	885.3(2.9) (x3.78)	C++; O1; chunk=96; batch=16
8 Python staged	pyAmpliCol - ASM	85.6 (x3.29)	1117.0(4.6) (x4.77)	1004.9(4.3) (x4.29)	ASM; O3; chunk=96; batch=16
8 Rusticol	pyAmpliCol - Rusticol	1.02e3 (x39.15)	487.7(9) (x2.08)	464.4(6) (x1.98)	Rusticol PyO3; C++ O3; chunk=96; cached NumPy input; reusable stage scratch; batch=128; two-stage save/load
9 reference	AmpliCol	39.4	N/A	567	Fortran AmpliCol; with -library=use
9 Python staged	pyAmpliCol - JIT	80.8 (x2.05)	6144(17) (x10.84)	6102(15) (x10.77)	JIT; batch=16; last successful JIT refresh
9 Python staged	pyAmpliCol - C++	2.83e3 (x71.80)	1504.9(11.7) (x2.66)	1252.7(9.6) (x2.21)	C++; O3; chunk=96; batch=128; chunk compile workers=10; two-stage save/load
9 Python staged	pyAmpliCol - ASM	349 (x8.85)	2688.2(5.5) (x4.74)	2459.1(4.5) (x4.34)	ASM; O3; chunk=64; batch=16
9 Rusticol	pyAmpliCol - Rusticol	2.83e3 (x71.80)	1164.8(9) (x2.06)	1120.9(1.0) (x1.98)	Rusticol PyO3; C++ O3; chunk=96; cached NumPy input; reusable stage scratch; batch=128; repackaged from existing C++ O3 artifact in 103 s; two-stage save/load

10 Why This Matches AmpliCol but Remains Python-Friendly

The Python layer owns graph construction, process metadata, deterministic phase-space steering, artifact generation, evaluator caching, and validation against `AMPLICOL`. It does not perform the hot current recursion component by component in production. The two production execution routes are now:

Python staged route: source fill $\rightarrow$ bounded stage evaluators $\rightarrow$ amplitude/raw-sum evaluator,

Rusticol route: one PyO3 call $\rightarrow$ Rust source fill and staged sweep $\rightarrow$ amplitude/raw-sum evaluator.

Thus `PYAMPLICOL` keeps `AMPLICOL`'s essential optimization - shared currents across helicity parents and a topological current-table sweep - while replacing generated Fortran subroutines by saved Spensio/Symbolica evaluator artifacts. The `rusticol` path removes the remaining Python orchestration cost by loading the same process directory and executing the sweep in Rust.

The old fully monolithic compiled-DAG route is now deprecated and kept only as reference code. Its alias-based design was useful for exploring whole-graph compilation, but the production target is the staged DAG: Python generation plus a Rust eager-DAG runtime over saved Symbolica stage evaluators. This keeps the runtime schedule close to `AMPLICOL`'s current-table library while leaving the graph compiler model-driven and process-generic.